UNIVERSITÄT BONN

Few-Shot Learning

Toward Few-Shot Learning and Data Augmentation

Bachelor Thesis

by

Milad Navidizadeh

2953248

in fulfillment of requirements for degree
Bachelor of Science (B.Sc.)

submitted to
Rheinische Friedrich-Wilhelms-Universität Bonn
Institute of Computer Science III
Bachlor Thesis for Information Systems and Artificial Intelligence

in degree course
Informatik (B.Sc.)

First Supervisor: Prof. Dr. Stefan Wrobel
                      University of Bonn

Second Supervisor: Prof. Dr. Christian Bauckhage
                      University of Bonn

Bonn, January 15, 2020

Acknowledgement

A special thanks goes out to Dr. rer. nat. Wolfgang Hansmann, for maintaining the guideline up to version 0.4 and Prof. Dr.-Ing. André Miede for his work on an extended template for `classicthesis` that I learned (and borrowed) a lot from.

> » *It has been a long journey to this moment.* «
> Sidney Poitier, 1964

Contents

1 Motivation

Nowadays machine learning and deep learning have become a distinguished approach for visual recognition tasks and has achieved great success in this process. However, they seek a large amount of labeled data to learn. Providing this amount of labeled data not only will bring much effort along but also will occupy a huge size of storage and seek large storage. In contrast, humans are very good in visual recognition so that, they can learn with one [1] or few [2] examples. Imagine one kid who can recognize a lion in a picture after looking a few pictures of lions as an example. We want to simulate and apply this human's ability to the deep learning and make them learn with few examples, with desirable accuracy. In this bachelor thesis, we concern ourselves with few-shot learning in deep learning. We aim to learn and train a model when there are few labeled examples obtainable. We approach to generate artificial examples from a few available labeled examples and enalrge our dataset artificially. This technique known as data augmentation. These artificial labeled examples aid us to learn better with more accuracy and prevent overfitting. Data augmentation is our focus in this work to achieve few-shot learning and prevent overfitting. In this thesis, we will introduce different well-known methods of data augmentation. The first purpose will be to discover if and how far data augmentation can improve the learning process and accuracy. The second step will be to compare their accuracy. In the end, we aim to discover the potential possibility of combining the different methods of data augmentation to increase accuracy and reduce error-rate and improve the learning process. We will focus on visual recognition tasks and their classification. Additionally, we will concentrate on the implementation of various methods of data augmentation for convolutional neural networks

---

[1] This is known as one-shot learning in deep learning and represents the scenario when there is one instance of each class in training-set to learn.

[2] This is known as few-shot learning in deep learning and represents the scenario when there are just few instances of each class in training-set to learn.

2 INTRODUCTION

Neural networks can possibly contain multiple non-linear layers and this makes them very expressive models that can learn very complicated relationships between their inputs and outputs. With even limited input data, neural networks can discover and learn many relations from the data, however, sometimes the discovered and learned relations do not exist or just consist of redundant information and relations. Redudant relation can potentially arise of data-noise or lack of data-generalization. Non exist relations can potentially emerge from lack of enough data. These phenomena known as *overffiting* in deep learning. In other words, learning with few labeled examples or noisy data causes overfitting. Overffiting cause low accuracy and high error-rate. Hence, we approach to propagation of artificial labeled examples from a few given examples to prevent overfitting and reduce the error rate and increase the accuracy.

As we mentioned aquiring a huge labeled dataset is expensive and seeks much effort and time. Therefore we aim to generate artificial example from few obtainable labeled examples. In other words, we augment our data and this strategy is known as data augmentation. There are a few well-known methods for data augmentation. We aim to introduce them in this thesis. Besides we will implement these methods to compare their efficiency and capability. These methods are as follow:

- **Label Preserving Transformation** 5.1
- **Elastic Distrotion** 5.2
- **Stroke Warping** 5.3
- **Bayesian Approach** 5.4

3 DATA REPRESENTATION

Here should be an Introdution of chapter and introduce the data

3.1 MNIST

The MNIST dataset (Modified National Institute of Standards and Technology) is a large handwritten digits dataset, provided by Yann Le Cun, derived from NIST Special Database 19 [STA].

The MSNIT dataset consists of $60,000$ train- and $10,000$ test-images and each image is grayscale with 28×28 pixels. It has 10 classes that represent $0 - 9$ digits and data is fairly splitted between classes [LeC]. MNIST is one of the most popular datasets for deep learning because of the not too high complexity and compatibility with almost all deep learning models. Hence many papers attempted to reach a low error-rate on this dataset. One of them manages to reduce the error-rate on the MNIST by up to 0.23% [CMS12]. You can find the information about the dataset at table 1 and figure 1 shows an example of the dataset.

FIGURE 1: *7 examples per class of MNIST dataset, merged in one image [Ath]*

Figure 2: *10 examples per class of CIFAR-10 dataset, merged in one image [Kri]*

Table 1: *Structure and organization of the datasets.*

Dataset	NO. Classes	NO. Train	NO. Test	Size (pixel)	NO. Channel
MNIST	10	60,000	10,000	28×28	1 (Grayscale)
Fashsion-MNIST	10				
CIFAR-10	10	50,000	10,000	32×32	3 (RGB)
CIFAR100	100				

3.2 Fashsion-MNIST

blablabla

3.3 CIFAR-10

The CIFAR-10 (Canadian Institute for Advanced Research) , collected by Alex Krizhevsky, Vinod Nair, and Geoffrey Hinton is a subset from 80 million tiny images dataset [Uni].

The dataset consists of $60,000$ RGB with 32×32 pixels images, which are divided to the $50,000$ train and $10,000$ test datasets. As the name makes it clear the CIFAR-10 contains 10 classes ([plane, car, bird, cat, deer, dog, frog, horse, ship, truck]) [Kri]. One of the lowest reported error-rate with a convolutional neural network managed to achieve 2.56% [DT17]. You can find the information about the dataset at table 1 and figure 2 shows an example of the dataset.

3.4 CIFAR100

blablabla

4 PRELIMINARIES

4.1 INTRODUCTION

4.2 CONVOLUTIONAL NEURAL NETWORK

4.3 OVERFFITING

5 Data Augmentaion

Here should be Introdution for data augmenation

5.1 Label Preserving Transformation

One of the most in common method to enlarge the dataset artificially is the label preserving transformation. This method provides the possibility to generate artificial data with non-heavy computation. The advantage of a very little computation aids us to save storage. In other words, it wouldn't be required to save the generated data on a storage and the data can always be enlarged artificially in a short time and with a little computation. As the method's name makes it clear, this method aims to generate new data from a single data with the same label. We explain the method and its approach for image datasets because as we mentioned our focus is image datasets.

This method consists of generating image translations and horizontal reflectins. Image translations mean, random patches from the original images. The size of the patches are smaller than the size of the original images. We will extract all pssobile translations (all possible patches which fit in the image) from our image. These extracted translations (extracted random patches) and their horizontal reflection will be used for training our network. Given a single channel (grayscale) n×n image and the translation pathce with size of m×m and $n > m$ then the training dataset will increase by a factor:

$$2 \times (n - m + 1) \times (n - m + 1)$$

In the figur 3 is the process of the translations and reflections of a small image presented.

At test time, the method extracts the patches with the same size. This time the patches will be extracted from 4 corners and center of the original image.

Figure 3: *An example of single channel image with size of* 4×4 *with its translations with size of* 3×3 *patches and their horizontal reflections. The numbers determinate the pixels intensity*

The network predicts on these five patches and their horizontal reflections (10 patches altogether). In the end, the average of the predictions will be our network's prediction.

5.2 Elastic Distrotion

Another well-known method for data augmentation is elastic distortion. This method as the same as label preserving transformation generates artificial data (images) from a single data (image) but with this difference that this time instead of resizing the image (translation) the pixels inside of the image will be moved and displaced. The intensity of pixels also will be changed with respect to their new places and their neighbors' intensity and their old places and their origin intensity.

As it cleared elastic distortions concerns itself with a displacement of the pixels and an adjustment of the pixels intensity. This approach is known as interpolation.

There are a few schemes like (nearest neighbor-, bicubic-, spline- and bilinear interpolation) to reach the goal. Bilinear interpolation is one of the simplest with a desirable result. Hence we use bilinear interpolation. We will explain the process below.

We will start with pixels displacement. For instance if $\Delta x(x,y) = \alpha x$ and $\Delta y(x,y) = \alpha y$, this means that each pixels will be moved αx in the direction of x-axis and αy in the direction of y-axis. α would be our scale parameter and since α can be a non-integer value, interpolation is necessary and as we mentioned we use the bilinear interpolation. In the next step after pixels displacement, the intensity of the pixels in the new locations should be adjusted concerning the intensity of the neighbors' pixels in the original image (origin square). Hence bilinear interpolation aids us to reach this target. The bilinear interpolation interpolates the moved pixel horizontally. Then it interpolates the pixel vertically with respect to yielded values from horizontal interpolations. We will show and summarize the process formally in below.

Definition 1: Given p' the pixel which we want to dispalacment it with $\Delta x(x,y) = \alpha x$ and $\Delta y(x,y) = \alpha y$ and $p_{(x,y)}, p_{(x+1,y)}, p_{(x,y-1)}, p_{(x+1,y-1)}$ are the neighbors (on origin square) of p' in the new location after displacment and $I(p)$ shows the intensity of pixel p. Then the vertical interploation yields:

$$V_1 = I\big(p_{(x,y)}\big) + \Big(\Delta x(p', p_{(x,y)}) \times I\big(p_{(x+1,y)}\big)\Big)$$
$$V_2 = I\big(p_{(x,y-1)}\big) + \Big(\Delta x(p', p_{(x,y-1)}) \times I\big(p_{(x+1,y-1)}\big)\Big)$$

And then the horizontal interploation yields a new intensity for pixel p' after displacment:

$$I(p') = V_1 + \Big(\Delta y(p', p_{(x,y-1)}) \times V_2\Big)$$

To reach elastic deformation or elastic distrotion we approach to generate $\Delta x(x,y) = \alpha x$ and $\Delta y(x,y) = \alpha y$ with $\alpha \times rand(-1,+1)$ since α is image-scale parameter and $rand(-1,+1)$ generate a number between -1 and $+1$ with uniform distribution. After all, the fields Δx and Δy are convolved with a Gaussian filter with standard deviation of σ. The values of σ and α depends on the image size and the image entropy. This process will generate elastic deformed image from orginal image which called elastic distrotion.

Figure 4: *An exmaple of rotation, skew, and shear (scale) transforamtions for stroke warping respectively from left to right [***TODO-Augmentaion***]*

5.3 Stroke Warping

This method as same as previously introduced methods uses predetermined families of transformations. In other words, we enlarge our dataset artificially with the aid of well-known classical computer vision transformations. This method notwithstanding of non-heavy complexity accomplished desirable results even on medical purposes [Bai+19]. The ideas of this methode raised from Tangent Dist [Sim+92] and Tangent Prop [SLD93] works.

In this method, we perform small changes in images to augment our data. That means we augment our images with skewing, rotating and shearing (scaling) them [YLW97]. As same as the label preserving transformation and elastic distortion the augmentation will be started before the training phase and the training will be performed on enlarged data. Figure 4 represents each mentioned transformation to make them visually understandable.

5.4 Bayesian Approach

As you may until now realized, all introduced classes of techniques in this work for data augmentation were pre-defined families of statistical transformations. In other words, we generated synthetic images with computer vision techniques statistically before the training phase and our augmented data generated once and during the learning process we didn't change them. In fact, the basic idea of each of them is known as "poor man's" data augmentation (PMDA) [Tan91]. The point is if neural networks, in general, can learn such complex relations and patterns in images, therefore they should be able to learn latent variables to improve data augmentation dynamically. Put the matter another way, in this method, we focus not only on the classification task and learning patterns of images but also on

learning how to improve our image augmentation dynamically and iteratively during the training and classification.

Before we start with the description of the Bayesian approach, it would be necessary to introduce generative models or more especially Generative Adversarial Networks (GANs). One of the remarkable works on GANs proposed by Goodfellow et al. [**TODO**]. In the below first, we will introduce GANs and after that and in the Bayesian approach description we will explain how this approach uses the main idea of GANs and extend it to improve data augmentation.

5.4.1 Generative Adversarial Network (GAN)

A Generative Adversarial Network (GAN), in general, consists of two parts:

- **Generator**: As the name cleared itself, it tries to generate new data and simultaneously learns how to generate the data more plausible.
- **Discriminator**: The discriminator learns to distinguish generated fake data from the generator and real data. Besides, it helps and teaches the generator like a kind teacher to generate more plausible data and close to real.

Indeed the generator and the discriminator play a minimax game. In other words, the generator tries to fool the discriminator with fake data and the discriminator tries to distinguish this matter and tell the generator that he failed to fool and helps the generator to generate data closer to real ones. In the beginning, the discriminator can simply distinguish the real data and the generator's fake data. But after each iteration, the generator becomes more intelligent and generates more plausible data. In the end, the generator becomes so intelligent so that it can almost fool the discriminator. This means, that the discriminator's accuracy reduces to about 50%. In effect, the discriminator predicts almost randomly while it predicts between 2 labels (real or fake). Sometimes it is even hard for humans to distinguish the generator's fake data and real data.

The generator will be fed by random[1] input to generate synthetic data. In general, GANs try to replicate a probability distribution. For this matter, they use loss function to measure the distance between the distribution of the generator's data and the distribution of the real data. As we mentioned Goodfellow el al. [**Same**] suggest minimax loss in their work. The minimax loss defined as:

$$Minimax\ Loss = E_x[\log(D(x))] + E_z[\log(1 - D(G(z)))] \qquad (5.1)$$

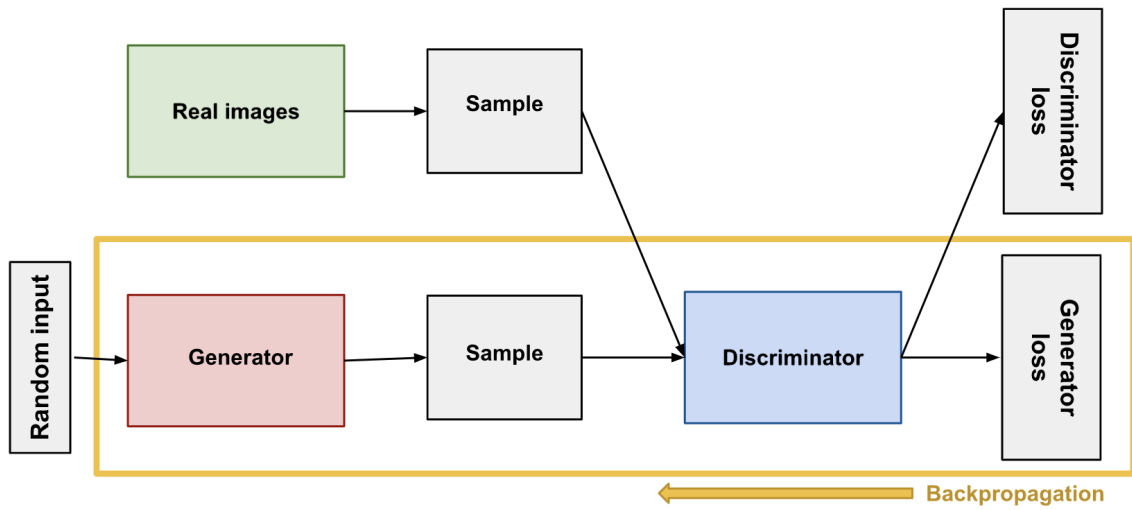[1] noise

Figure 5: *GAN architecture [***TODO-google-reference***]*

where $D(x)$ denotes the discriminator's estimate of the probability that real data instance x is real and $D(G(z))$ denotes the discriminator's estimate of the probability that a fake instance is real. It is clear that the Generator tries to minimize the (5.1) and the discriminator tries to maximize it. Figure 5 shows a basic architecture of a GAN.

Now and after briefly introduction of GANs we can start with the **Bayesian Approach**. One of the noteworthy work for data augmentation with the aid of the bayesian model proposed by Toan Tran et al. [Tra+17]. In this approach, our deep learning model tries to estimate the distributions of labeled data and with aid of the estimated distributions optimizes the latent variable used for data augmentation. The approach uses the GANs architecture skeleton to generate synthetic data with the difference that the optimization of latent derived from the Bayesian model. The principle idea for data augmentation using latent variables proposed by the statistical learning community [TW87]. Nevertheless applying the idea, directly into deep learning seeks a massive computational effort. Therefore we talked before about the estimation. To be more precise, the approach uses a novel Bayesian data augmenation algorithm, called Generalized Monte Carlo Expectation Maximization (GMCEM). This algorithm augments training data and mutually optimizes the network parameters. The algorithm successively generates synthetic data and use Monte Carlo to estimate the expected value of the network parameters given the previous estimate instead of calculating loss function. After the estimation of the expected value, the parameter values will be updated with stochastic gradient descent (SGD). In the end, the algorithm and approach turned in to reality with the aid of GANs. The proposed GAN consists of one generator

Figure 6: *The network architecture of Bayesian data augmentaion approch [Tra+17]. G: Generator, A: Authenticator, C: Classifier.*

and 2 discriminators. As we in the GANs section (5.4.1) discussed the generator is responsible to generate synthetic data and one of our discriminators distinguishes fake and real data. However, the second discriminator discriminates between the classes of data. Figure 6 represents the utilized netwok architecture in this approach visually. This proposed architecture is nearly similar to the Auxiliary classifier GANs (AC-GANs) [Aa16]. Nevertheless, in the AC-GANs discriminator responsible for both classification real-or-fake data and data labels (classes) and in our network we utilized 2 discriminators separately for this matter.

In the following, we will explain the utilized algorithm formally from the Toan Tran el al. work [Tra+17]. As we mentioned the goal is to estimate the parameters of the neural networks using labeled data. The training process is defined by the following optimization problem:

$$\theta^* = \arg\max \log p(\theta|\mathbf{y}) \tag{5.2}$$

Where training set denoted as $\mathcal{Y} = \{\mathbf{y}_n\}_{n=1}^N$ with $y = (t, x)$ and $t \in \{1, ..., K\}$ (Classes-Set) and data samples $\mathbb{R}^D$ and θ denoted as model (network) parameters and observed posterior defined as:

$$p(\theta|\mathbf{y}) = p(\theta|t, \mathbf{x}) \propto p(t|\mathbf{x}, \theta) p(\mathbf{x}|\theta) p(\theta) \tag{5.3}$$

Now if we assume that the data samples $\mathcal{Y}$ are conditionaly independent, we can define the following loss function which maximize the (5.2).

$$\log p(\theta|\mathbf{y}) \approx \log p(\theta) + \frac{1}{N}\sum_{l}^{N}\left(\log p\left(t_n|\mathbf{x}_n,\theta\right) + \log p\left(\mathbf{x}_n|\theta\right)\right) \tag{5.4}$$

where $p(\theta)$ denotes a prior on the distribution of the deep learning model parameters, $p(t_n|x_n,\theta)$ represents the conditional likelihood of label t_n, and $p(x_n|\theta)$ is the likelihood of the data x.

After estimation and optimization the θ on our training set, it is the time to generate synthetic data from y using latent variable z. Therefore the augmented $p(\theta|y,z)$ can be estimated. The latent variable z as same as y defined as $z = (t^\alpha, x^\alpha)$ where $t^\alpha \in \{1,...,K\}$ denotes associated label and $x^\alpha \in \mathbb{R}^D$ is synthesized sample. As we mentioned to avoid a heavy and most likely interminable computation, instead of the Expectation-Maximization (EM) algorithm we will use Generalized Monte Carlo EM Algorithm to estimate the expected value and maximize it.Hence the augmented posterior $p(\theta|y,z)$ for latent variable z will be as follow:

$$p(\theta|\mathbf{y},\mathbf{z}) = \frac{p(\mathbf{y},\mathbf{z},\theta)}{p(\mathbf{y},\mathbf{z})} = \frac{p(\mathbf{z}|\mathbf{y},\theta)p(\theta|\mathbf{y})p(\mathbf{y})}{p(\mathbf{z}|\mathbf{y})p(\mathbf{y})} = \frac{p(\mathbf{z}|\mathbf{y},\theta)p(\theta|\mathbf{y})}{p(\mathbf{z}|\mathbf{y})} \tag{5.5}$$

where the expectation step will be defined as follow:

$$p(\theta|\mathbf{y},\mathbf{z}) = \frac{p(\mathbf{y},\mathbf{z},\theta)}{p(\mathbf{y},\mathbf{z})} = \frac{p(\mathbf{z}|\mathbf{y},\theta)p(\theta|\mathbf{y})p(\mathbf{y})}{p(\mathbf{z}|\mathbf{y})p(\mathbf{y})} = \frac{p(\mathbf{z}|\mathbf{y},\theta)p(\theta|\mathbf{y})}{p(\mathbf{z}|\mathbf{y})} \tag{5.6}$$

where $\mathbf{z}_m \sim p\left(\mathbf{z}|\mathbf{y},\theta^i\right)$, for $m \in \{1,...,M\}$. In (6), if the label t_m^a of the m^{th} synthesized sample $\mathbf{z_m}$ is known, then $\mathbf{x}_m^a$ can be sampled from the distribution $p\left(\mathbf{x}_m^a|\theta,\mathbf{y},t_m^a\right)$. Hence, the conditional distribution $p(\mathbf{z}|\mathbf{y},\theta)$ can be decomposed as:

$$p(\mathbf{z}|\mathbf{y},\theta) = p\left(t^a,\mathbf{x}^a|\mathbf{y},\theta\right) = p\left(t^a|\mathbf{x}^a,\mathbf{y},\theta\right)p\left(\mathbf{x}^a|\mathbf{y},\theta\right) \tag{5.7}$$

where $(t^a,\mathbf{x}^a)$ are conditionally independent of y given that all the information from the training set y is summarized in θ this means that $p\left(t^a|\mathbf{x}^a,\mathbf{y},\theta\right) = p\left(t^a|\mathbf{x}^a,\theta\right)$, and $p\left(\mathbf{x}^a|\mathbf{y},\theta\right) = p\left(\mathbf{x}^a|\theta\right)$. Now with respect to the θ for the maximization step with concern of removing the independent terms for θ will derived the maximization of $\hat{Q}\left(\theta,\theta^i\right)$ as follow:

$$\hat{Q}\left(\theta,\theta^i\right) = \log p(\theta) + \frac{1}{N} \sum_{n=1}^{N} \left(\log p\left(t_n|\mathbf{x}_n,\theta\right) + \log p\left(\mathbf{x}_n|\theta\right)\right) + \frac{1}{M} \sum_{m=1}^{m} \log p\left(\mathbf{z}_m|\mathbf{y},\theta\right)$$

$$=$$

$$\log p(\theta) + \frac{1}{N} \sum_{n=1}^{N} \left(\log p\left(t_n|\mathbf{x}_n,\theta\right) + \log p\left(\mathbf{x}_n|\theta\right)\right) + \frac{1}{M} \sum_{m=1}^{n} \left(\log p\left(t_m^a|\mathbf{x}_m^a,\theta\right) + \log p\left(\mathbf{x}_m^a|\theta\right)\right)$$

$$(5.8)$$

After all, we estimate the θ^{i+1} so that $\hat{Q}(\theta^{i+1},\theta^i) > \hat{Q}(\theta^i,\theta^i)$. To reduce the computation complexity as we mentioned instead of gradient descent, stochastic gradient decent (SGD) utilized for estimating the θ^{i+1}. The iteration will be continued until $|\theta^{i+1} - \theta^i|$ get sufficiently small.

As we made it clear the above formal explanations and equations are derived from [Tra+17] and in some points are matched one-to-one.

6 Expriments & Results

6.1 CNNs Architecture

In this section, we will introduce our classifiers (CNNs) architecture for various datasets which we introduced in the chapter (3). We picked our CNNs architecture from different sources regarding their non-heavy complexity and desirable accuracy. Maybe our chosen CNNs architecture doesn't provide the best-reported accuracies on datasets, but while we use the same CNN-architecture for each dataset and we aim to compare various data augmentation approaches on each dataset, it would not affect our results. Moreover, as we will report in section (6.3) our accuracies are not so far away from best-reported accuracies.

6.1.1 MNIST

For the MNIST dataset, we have chosen a semi-simple CNN architecture from [rep17] with 2 convolutional layers and 2 fully connected layers. ReLU function is utilized as the activation function. For training and to calculate the loss function, **CrossEntropyLoss ??** and to optimize the network's parameters **Adam optimizer ??** from the Pytorch have been chosen. The learning rate for Adam optimizer is set to 0.001. To reduce overfitting, a drop-out layer is placed at the end of the second convolutional layer. Figure 7 represents the explained CNN architecture visually.

6.1.2 Fashsion-MNIST

6.1.3 CIFAR-10

For the CIFAR-10 dataset, we have chosen a bit more complex CNN architecture from [Zhe18] with 6 convolutional layers and 3 fully connected layers. ReLU function is utilized as the activation function. For training and to calculate the loss function, **CrossEntropyLoss** and to optimize the network's parameters **Adam optimizer** as same as 2 previous CNNs architecture have been chosen. The learning

FIGURE 7: *CNN Architecture for training the CIFAR-10 dataset [17]*

rate for Adam optimizer is set again to 0.001. To reduce overfitting, a drop-out layer is placed at the end of the fourth convolutional layer. Figure 8 represents the explained CNN architecture visually.

6.2 Implementaions

6.2.1 Label Preserving Transformations

6.2.2 Elastic Distortion
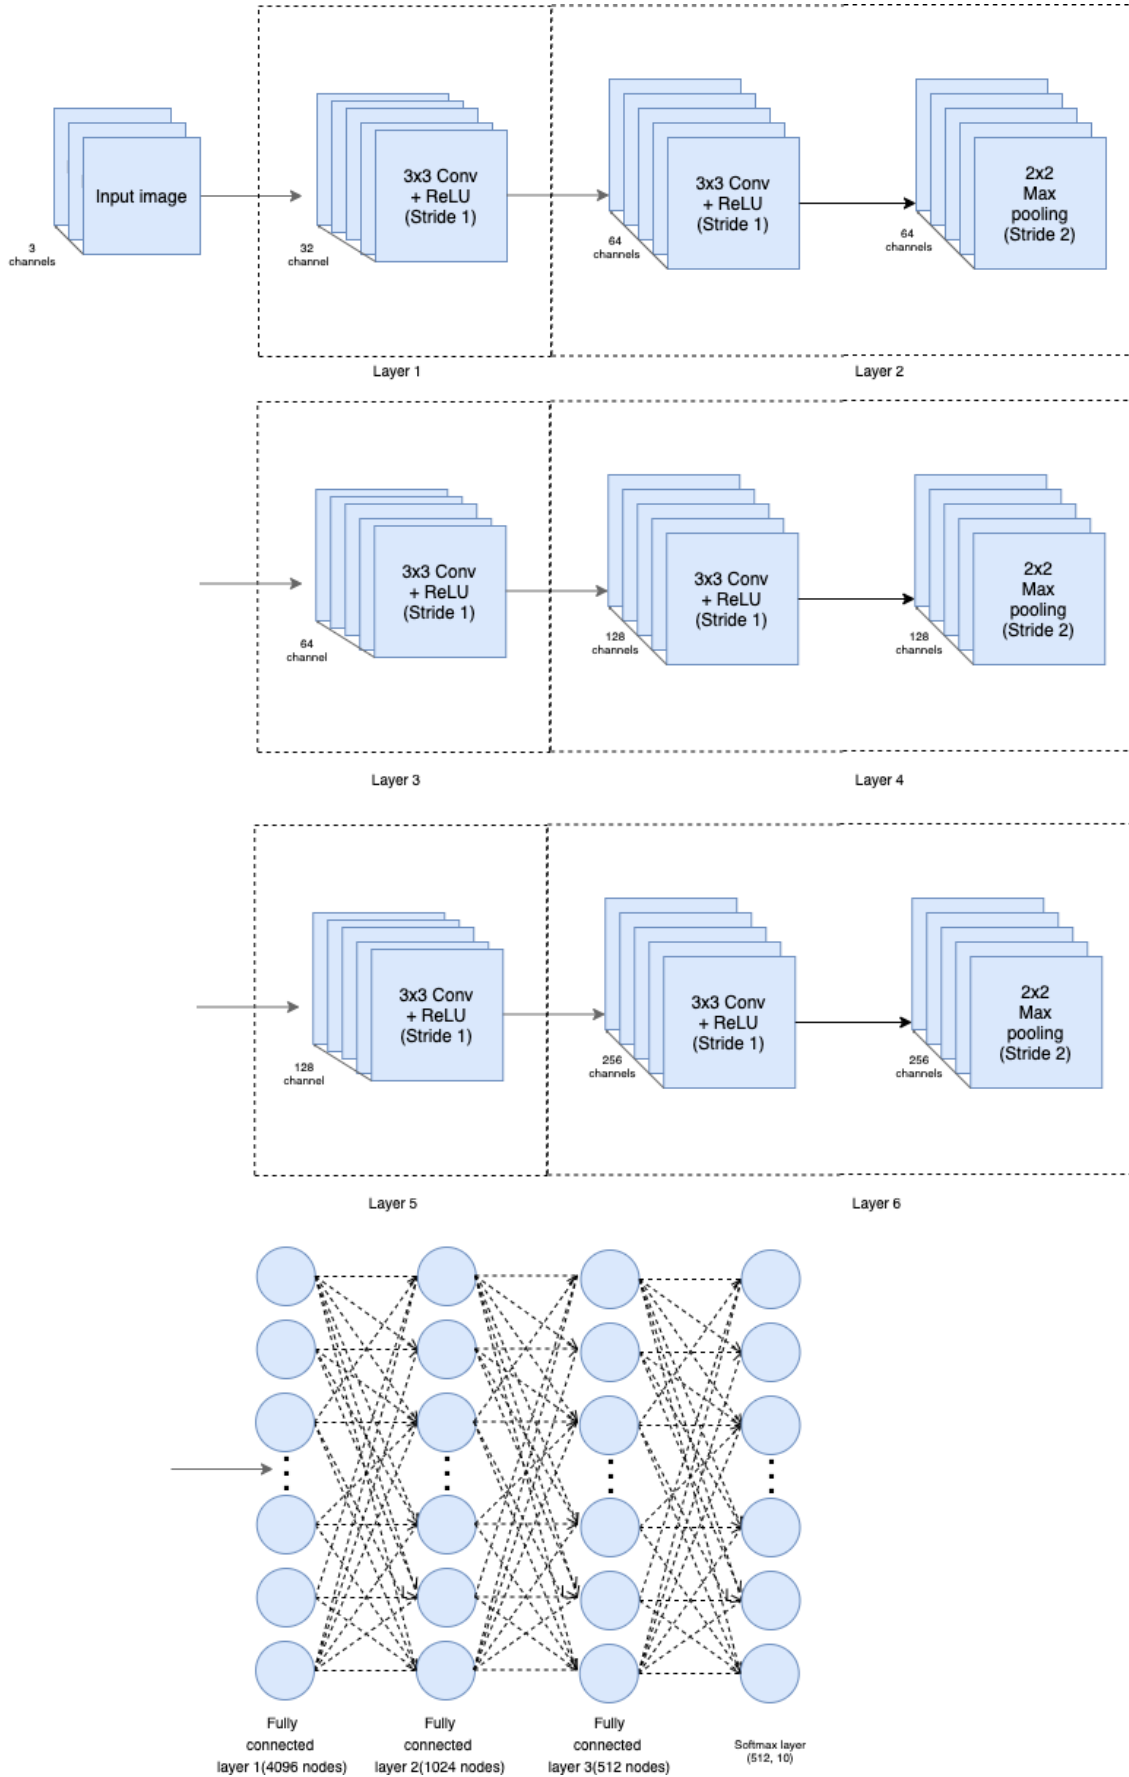
6.2.3 Stroke Warping

6.2.4 Bayesian Approach

6.3 Results

Figure 8: *CNN Architecture for training the CIFAR-10 dataset*

7 Comprehensive Comparison

8 Contribution Of Work

9 BIBLIOGRAPHY

[17]   *Convolutional Neural Networks Tutorial in PyTorch.* https : / / adventuresinmachinelearning . com / convolutional – neural – networks – tutorial – in-pytorch/. Accessed: 2020-01-15. 2017.

[Aa16]   A.Odena and C.Olah andJ.Shlens. *Conditional Image Synthesis With Auxiliary Classifier GANs.* 2016. arXiv: 1610.09585 [cs.CV].

[Ath]   P. Athul. *Medium Handwritten digit recognition using PyTorch.* https://medium.com/ @athul929/hand-written-digit-classifier-in-pytorch-42a53e92b63e. Accessed: 2019-12-16.

[Bai+19] Sung W. Baik et al. "Multi-grade brain tumor classification using deep CNN with extensive data augmentation". In: 30 (Jan. 2019). https://www.sciencedirect.com/ science/article/pii/S1877750318307385, pp. 174–182.

[CMS12] Dan Cireşan, Ueli Meier, and Juergen Schmidhuber. *Multi-column Deep Neural Networks for Image Classification.* 2012. arXiv: 1202.2745 [cs.CV].

[DT17]   Terrance DeVries and Graham W. Taylor. *Improved Regularization of Convolutional Neural Networks with Cutout.* 2017. arXiv: 1708.04552 [cs.CV].

[Kri]   Alex Krizhevsky. *The CIFAR-10 dataset (Canadian Institute for Advanced Research).* http: //www.cs.toronto.edu/~kriz/cifar.html. Accessed: 2019-12-16.

[LeC]   Yann LeCun. *exdb THE MNIST DATABASE of handwritten digits.* http://yann.lecun. com/exdb/mnist/. Accessed: 2019-12-16.

[rep17] GitHub repository. *adventures-in-ml-code.* https : / / github . com / adventuresinML / adventures-in-ml-code. Accessed: 2020-01-15. 2017.

[Sim+92] P. Simard et al. "Tangent Prop-A Formalism for Specifying Selected Invariances in an Adaptive Network". In: *Advances in Neural Information Processing Systems 4.* Morgan Kaufmann, 1992, pp. 895–903.

[SLD93] P. Simard, Y. LeCun, and .T. Denker. "Efficient Pattern Recognition Using a New Transformation Distance". In: *Advances in Neural Information Processing Systems 5.* Morgan Kaufmann, 1993, pp. 50–58.

[STA]   STANDFORD. *NIST National Institute of Standards and Technology.* https://www.nist. gov/data. Accessed: 2019-12-16.

[Tan91] Martin A. Tanner. *Tools for Statistical Inference. Observed Data and Data Augmentation Methods.* 1st ed. Vol. 67. ISBN 978-0-387-97525-2. Springer-Verlag New York, 1991.

[Tra+17] T. Tran et al. *A bayesian data augmentation approach for learning deep models*, 2017. arXiv: 1710.10564 [cs.CV].

[TW87] Martin A. Tanner and Wing Hung Wong. *The Calculation of Posterior Distributions by Data Augmentation. Journal of the American Statistical Association.* 1st ed. Vol. 82. ISBN 82(398):528–540. American Statistical Association, 1987, pp. 528–540.

[Uni] New York University. *Visual Dictionary Teaching computers to recognize objects.* http://groups.csail.mit.edu/vision/TinyImages/. Accessed: 2019-12-16.

[YLW97] Larry S. Yaeger, Richard F. Lyon, and Brandyn J. Webb. "Effective Training of a Neural Network Character Classifier for Word Recognition". In: *Advances in Neural Information Processing Systems 9.* Ed. by M. C. Mozer, M. I. Jordan, and T. Petsche. MIT Press, 1997, pp. 807–816. URL: http://papers.nips.cc/paper/1250-effective-training-of-a-neural-network-character-classifier-for-word-recognition.pdf.

[Zhe18] Zhenye. *Deep Learning with Pytorch on CIFAR10 Dataset.* https://zhenye-na.github.io/2018/09/28/pytorch-cnn-cifar10.html. Accessed: 2020-01-15. 2018.

*

List of Figures

List of Tables

Statement of Authorship

I hereby confirm that the work presented in this bachelor thesis has been performed and interpreted solely by myself except where explicitly identified to the contrary. I declare that I have used no other sources and aids other than those indicated. This work has not been submitted elsewhere in any other form for the fulfilment of any other degree or qualification.

Bonn, January 15, 2020

_____

Milad Navidizadeh